

Half of Your CMDB's Measured Value Is Knowing Who Opened the Ticket

Sudhir Dixit

Independent Researcher
sudhir.dixit1@gmail.com

Abstract

Configuration management database (CMDB) programmes are justified by the analytics they enable, and that justification is a comparison against a baseline the analyst chooses. On a public event log of 45,455 incidents from a bank's IT operation whose affected-item field is 100% populated, we show the choice is worth as much as the data. Predicting at intake whether an incident will be reassigned, knowing the affected configuration item is worth +0.183 AUC against four intake fields, and +0.103 [+0.094, +0.113] once one more free field is added: which group logged the incident. The gap survives three estimator families, including one in which item identity is a single column rather than 2,554, and reproduces on two further targets and at three target thresholds. Stated as a detection problem it is larger than the AUC ratio of 1.8 suggests — at a review capacity of 5% of arrivals, omitting the field credits the CMDB with 3.9 [3.2, 6.4] times as many additional catches. Two measurements explain the overlap: the opening group adds under 0.01 AUC to a model that already knows the item, and the shrinkage is graded in the field's resolution, running 27% to 44% as it goes from one bit to 49 levels. We also report three things we got wrong: a null that could not fail, an operational factor we first reported without an interval, and a field whose meaning we mis-stated in an earlier draft.

1 Introduction

A CMDB programme is expensive and is justified by what it will enable. That justification rests on a comparison whose second arm — performance without configuration data — is a choice the analyst makes, usually silently, by deciding which already-available fields to include.

On one organisation's incident log the choice halves the answer, and we can say exactly where the difference goes. This is a short paper with one result and one mechanism.

What kind of contribution this is. We report no deployed system. This is a retrospective study on a public benchmark, and it belongs to this track as an evaluation practice: the quantity we measure is not a model's accuracy but the number a deployment decision turns on, and we show that number is set by a field-admission choice that is rarely written down. The organisations that make this choice are making it now,

on data of exactly this shape, and Section 6 states what it costs them in the units a service desk operates in.

On what this paper used to claim. Earlier versions reported larger findings withdrawn under adversarial review, each an artifact of a control we built ourselves. Three corrections survive here as reported results rather than silent edits, because each is the kind of error this paper is about: a mechanism null that could not fail (Section 5), an operational factor first reported without an interval (Section 6), and — the one that cost us a title — a central field whose meaning we asserted rather than checked (Section 3).

2 Background

The CMDB and what justifies it. A CMDB records configuration items and their relationships, and is the asset on which the configuration-management practice of ITIL and its successors rests. Survey work on ITSM framework adoption (Marrone and Kolbe 2011) finds that realised benefits grow with implementation maturity, but measures them as practitioner-reported outcomes rather than as the analytic value of the configuration data itself. We are not aware of a peer-reviewed measurement of what knowing the affected configuration item is worth to a named prediction task; the material addressing that question directly is vendor documentation. The gap is the occasion for this paper, and part of the reason it exists is that the data to close it barely does: of the three public ITSM incident logs in common use, one records the affected item on every incident, the second records it on 0.2%, and the third has no such field at all (van Dongen 2014; Amaral, Fantinato, and Peres 2018; Steeman 2013). That is also why this study is single-organisation, and it is a constraint rather than a choice.

The methods a CMDB is bought to enable. Surveys of AIOps for failure management (Zhang et al. 2024; Remil et al. 2024; Notaro, Cardoso, and Gerndt 2021) catalogue triage, failure-prediction and root-cause methods that take configuration data as input. Incident triage in particular — deciding which team owns an incident — has been treated as a learning problem on operational logs at industrial scale (Chen et al. 2019), as has failure prediction across a cloud estate (Lin et al. 2018).

Prediction on incident event logs. Our task sits in predictive process monitoring, where outcome prediction is surveyed and benchmarked across real-world logs (Teinemaa et al. 2019; Verenich et al. 2019) on data of the kind process mining formalises (van der Aalst 2016). The BPI Challenge logs are the standard public instances (van Dongen 2014; Steeman 2013) and comparable ITSM exports exist elsewhere (Amaral, Fantinato, and Peres 2018). Such logs carry documented quality hazards: the imperfection patterns catalogued by Suriadi et al. (2017) include several we rule out explicitly below, in particular values denormalised from a record onto its events.

What the comparison is measured against. The quantity we report — the value of a feature *given* a baseline — is what algorithm-agnostic variable-importance frameworks formalise (Williamson et al. 2021), and the overlap we measure is the feature interaction additive importance measures are built to handle (Covert, Lundberg, and Lee 2020). That measured importance depends on what the model already has is therefore not news as a statistical fact; what is not established is its magnitude for configuration data on a named operational task, or that the field it overlaps with is one the organisation already records for nothing. Rendle, Zhang, and Koren (2019) show for recommender systems that inadequately specified baselines can account for reported gains, and Ferrari Dacrema, Cremonesi, and Jannach (2019) report the same across neural recommendation work. Benchmark conclusions are sensitive to variation that is seldom reported (Bouthillier et al. 2021), and documentation of the choices producing a result is thin (Gundersen and Kjensmo 2018). Where the choice concerns which fields may enter a model it shades into leakage (Kaufman et al. 2012; Kapoor and Narayanan 2023); for predictive process monitoring, Weytjens and De Weerd (2022) show benchmark datasets need explicit construction to avoid it. We report this failure mode for a deployment decision rather than a leaderboard, where field admission is not a modelling detail but the number the business case turns on.

3 Data and Task

We use the incident records of the BPI Challenge 2014 log (van Dongen 2014) from Rabobank Group ICT, recorded in HP Service Manager, with the activity log shipped alongside.

Cleaning. Of 46,809 rows, 203 carry no incident identifier and one an unparseable timestamp. A further 1,150 incidents opened before 2013-10-01 are left-censored — already open when the extract began — and show reassignment rates of 76–100% against a subsequent 35–42% that declines gently across the window. Removing them leaves 45,455 incidents to 2014-03-31; a temporal 70/30 split gives 31,818 training and 13,637 test incidents, positive rates 0.413 and 0.372. The cutoff is not taken from the outcome: monthly volume rises from 857 incidents in September 2013 to 8,606 in October. Over the analysed rows the affected-item field is 100% populated across 2,929 items, 2,554 of them seen in training.

Task and estimator. Predict at creation whether an incident will later be reassigned. We call the target reassign-

ment rather than misrouting throughout: it fires on 37.2% of test incidents, which is routine handling and not error, and Section 6 reports the ladder at stricter thresholds. Median handling time rises from 1.59 hours for incidents never reassigned to 47.53 hours for those reassigned five or more times; we use this to motivate the task, not as a recoverable saving. The primary estimator is one-hot logistic regression, chosen because it is bin-free; two further estimator families are reported in Section 4. Encoders and rankings are fit on training data only; intervals are 2,000-resample paired bootstraps.

The free field, and what it is not. The activity log records, on the `Open` event of every incident, an assignment group: 49 in training, none missing. An earlier draft of this paper called it the queue the service desk routed the incident to. That was an assertion, not a finding, and it is wrong. Three checks say so. The `Open` rows carry 50 distinct groups where the `Assignment` rows carry 218, and a routing destination cannot be less diverse than the teams that then do the work. One group accounts for 67.0% of `Open` rows but only 18.4% of all activity rows. And the `Open` group matches the group on the incident’s first `Assignment` activity for just 15.1% of incidents. The field is the group that *logged* the incident — a central service desk for two thirds of tickets, a specialist team opening its own work for the rest. It is also concentrated and unstable: the largest group holds 62.1% of training incidents, and the number of live groups falls from 49 to 32 between our two halves.

We report the error because the correction carries the paper’s method. Nothing measured changes — the field is still on the intake event, still 100% populated, still free, and still absorbs 44% of the item column’s value. What changed is a claim about what the field *meant*, made on no evidence, in the way this paper accuses business cases of treating baselines.

Why not the real routing decision. The obvious repair — use the first `Assignment` group instead — is not available. That activity occurs strictly after the `Open` event on every incident that has one, a median of 46 minutes later, and 7,878 of 45,455 incidents never have one at all. It fails the admissibility criterion below, and substituting it would introduce exactly the leakage that criterion exists to prevent. The opening group is admissible precisely because it is a property of the intake event itself.

Admissibility, and mutation. A field on the `Open` row is admitted only if it is a per-event observation, tested by whether it varies across an incident’s activity rows; the opening group varies for 92.56% of cohort incidents and matches the last-observed group for only 21.35%. Because the incident file is a closed-record snapshot we also report how often each field is edited after creation: Urgency on 1,107 incidents (2.44%), Impact on 1,084 (2.38%), the affected item on 159 (0.35%). Incidents whose Impact or Urgency was edited are more likely to be reassigned (0.66 against 0.39), so this is not neutral; Section 4 reports the headline with those fields dropped and with the edited incidents removed.




Figure 1: (a) The same quantity under three baselines; the third is discussed in Section 8. (b) The ordering holds at every temporal split point from 55% to 80%.

Baseline	AUC	+ item id.	gain [95% CI]
intake only	0.562	0.746	+0.183 [+0.172, +0.195]
+ opening group	0.644	0.748	+0.103 [+0.094, +0.113]

Table 1: The value of knowing the affected configuration item, with and without a field the service desk already records.

4 The Result

Against four intake fields the affected item looks decisive. Adding the opening group cuts its measured value by 44%. The gains are 28.1 and 17.4 standard deviations outside a null in which the incident-to-item association is shuffled, preserving column count and marginal distribution while destroying the signal — a comparison that matters because adding item identity adds 2,554 indicator columns, which is not free under a fixed penalty. These z figures pool the null’s Monte-Carlo dispersion with the gain’s own bootstrap standard error; dividing by the null alone, as an earlier draft did, would report 51 and 33.

The result is not a split artifact: across six temporal split points the two gains range +0.175 to +0.193 and +0.091 to +0.118 (Figure 1b). The bootstrap interval conditions on every design choice, so it is narrower than the honest spread: across split point, target threshold and cleaning cutoff the second gain ranges +0.068 to +0.130. Nor does the disclosed field editing explain it: reducing the intake block to Category alone — dropping Impact and Urgency, which are edited, and Priority, which is a deterministic function of them — gives +0.106, and restricting to the 44,227 never-edited incidents

gives +0.107.

It is not an artifact of the estimator. The dimensionality null above rules out the regularisation burden of 2,554 sparse columns, but a reader may reasonably want the effect shown under an estimator where that burden cannot arise at all. We therefore re-represent item identity as a single cross-fitted target-encoded column — fitted on training only, out-of-fold within training, so no row sees its own outcome — which makes the item one continuous feature rather than thousands of indicators, and makes histogram boosting usable on a variable that otherwise collapses 2,554 items into 137 bins. Across one-hot logistic regression, logistic regression on the encoded item, and boosting on the encoded item, the first rung ranges +0.173 to +0.183, the second +0.092 to +0.103, and the shrinkage 43% to 47%. Encoding a *shuffled* item column under the same pipeline returns -0.0001 ± 0.0016 and $+0.0042 \pm 0.0020$, so the encoder is not leaking; we report the second figure rather than round it away, and it bounds how much of the boosting rung could be encoding artifact.

5 Where the Difference Goes

The two rows differ by 0.080. That number is not itself informative — for any two variables, the drop in one’s measured gain when the other joins the baseline equals the drop in the other’s, so it restates the table rather than explaining it. Two direct measurements do explain it.

Read four ways, the overlap is this: the item is worth +0.183 over intake and +0.103 once the group is present; the group is worth +0.082 over intake and +0.002 once the item is present.

First: adding the opening group to a model that already knows the affected item is worth $+0.0017$ [$+0.0001, +0.0034$]. It is positive, and it clears a matched-dimension null of -0.0009 ± 0.0005 that no draw reaches — so it is not the cost of adding columns. But across split points and penalties it ranges $+0.0002$ to $+0.0072$, so the honest statement is that the group’s unique contribution is *under 0.01 AUC and not resolvable more finely*. Whatever it contributes, the item column has almost entirely absorbed it.

Second: the group’s predictive content is largely a function of the item. Randomising the group label *within* each item — which destroys the label’s own identity while preserving what the item implies about it — retains 91% of the group’s $+0.082$ gain.

A floor that could not fail. An earlier draft compared that 91% against a floor of 2% and claimed a margin of 89 points. The floor was wrong. It drew cell labels independently *per row*, so two incidents sharing an item landed in different cells and the association was destroyed by construction: the null could only ever return zero. Rebuilt as a random partition *of items*, which preserves item-level grouping and knows nothing about who opened anything, the floor depends on how fine the partition is — 41% at 49 cells, matching the field’s own cardinality, and 77% at 400. The real leg sits above every floor we can construct, but the honest margin at matched cardinality is 50 points, not 89, and the quantity is a function of a knob rather than a constant. We report the sweep (Figure 3a) and claim only the ordering.

How far this goes, and where it stops. Without any estimator, item identity carries 60.4% of the opening group’s information and the group carries 19.6% of the item’s — real, and strongly asymmetric, which is why we do not run the argument backwards. But it is not determination: a modal-group lookup keyed on item identity reproduces which group logged the incident on 90.3% of test incidents against 78.6% for always guessing the largest group, and class-balanced it reaches only 34.1%. Item identity resolves the coarse contrast — was this logged by the central desk or by a specialist team — and much less of the fine one.

The shrinkage is graded in the field’s resolution. That reading makes a prediction, and it holds. Replacing the group with progressively coarser versions of itself and re-running the same ladder gives shrinkages of 27%, 28%, 36% and 44% at two, four, eleven and 49 levels (Figure 3b). A single binary flag — central desk or not — already recovers 61% of the group’s baseline gain and 63% of the shrinkage it causes. The field that has to be admitted is close to one bit, which is far cheaper to reproduce in another organisation than a 49-way taxonomy.

What we do not claim. The reverse framing — that the item column “stands in for” the group — is not established. We measured it, by randomising items within groups, and obtained 44% of the item’s gain; but a random 50-cell grouping of items that knows nothing about intake retains 25%, and a grouping matched on the group’s mass profile retains -7% . Defensible nulls disagree about the floor, so the number is

Capacity	honest extra	naive extra	factor
5%	+67 [43, 84]	+262 [242, 300]	3.9 [3.2, 6.4]
10%	+33 [6, 70]	+353 [304, 404]	10.7 [5.5, 39.9]
20%	+18 [−28, 82]	+481 [407, 524]	26.7 [5.8, 221.0]

Table 2: Extra reassignment-bound incidents surfaced by adding item identity, against a baseline that includes the opening group (honest) and one that omits it (naive). Paired bootstrap over test rows, with ties at the capacity cut broken at random in every draw.

not interpretable and we exclude it. Nor do we claim a direction of causation: the log records no counterfactual, so whether the item drives who logs the ticket or both follow a third cause is not identified here. Our claim is about what an analyst’s baseline choice does to a measured number, which does not depend on which is true.

6 What the Difference Buys

An AUC difference is not a quantity a service desk can act on, so we state the same comparison as a detection problem. Fix a review capacity — the share of arrivals a desk can afford to look at before routing — rank the test incidents by predicted risk, and count how many reassignment-bound incidents each model surfaces. Both arms use matched models: the honest contrast adds item identity to intake-plus-group, the naive one adds it to intake alone, mirroring the two rows of Table 1. An earlier draft did not match them, crediting item identity with the group’s own contribution and inflating the factor by about a fifth.

At 5% capacity, knowing the affected item surfaces 67 [43, 84] more reassignment-bound incidents than the intake-plus-group baseline does, or about 36 a month at this log’s volume. An analyst who had omitted the group would have credited item identity with 262 [242, 300]. The omission overstates the operational gain by a factor of 3.9 [3.2, 6.4], against a ratio of only 1.8 between the same two gains measured as AUC.

The factor grows with capacity — 10.7 at 10%, 26.7 at 20% — but so does its interval, because the honest denominator shrinks to a few dozen incidents and at 20% its interval includes zero (Table 2). We therefore quote the 5% figure, which is the best resolved, and claim only that the overstatement is several-fold and larger than the AUC ratio implies. The reason it is larger at all is structural: AUC averages over every operating point, while a capacity-limited desk works at one, near the top of the ranking where the group-aware baseline is already close to its ceiling. What none of this says is that a surfaced incident is a corrected one — the log records no intervention, so the recoverable share is not identified here.

Three target definitions. Because reassignment is common, we repeat the ladder at stricter thresholds. Requiring two or more reassignments (21.8% of incidents) gives $+0.131$ and $+0.068$; three or more (10.6%) gives $+0.151$ and $+0.080$. The magnitude is threshold-dependent and not monotone, but every interval excludes zero, the ordering




Figure 2: (a) The overlap, drawn to scale: each circle’s area is that field’s measured gain over the intake block, and the centre distance is solved so that the intersection equals the measured overlap. The opening group lies almost entirely inside the item. (b) Randomising the group label within each item costs little; the same randomisation within a random partition of items costs more, by an amount that depends on how fine the partition is.

never reverses, and the shrinkage stays between 44% and 48%. The transferable quantity is the shrinkage, not the gain.

7 Does It Survive a Change of Task?

A gap on one target may be a property of that target, so we repeated the ladder on two further outcomes from the same cohort, changing only what is predicted. *Reopening* is close to an independent failure mode: it fires on 2,096 incidents and correlates with reassignment at +0.14. *Long handling* — handle time above the training third quartile — correlates at +0.40, so we report it as corroboration rather than as independent evidence.

The shrinkage reproduces on both. On reopening, item identity is worth +0.083 against the intake block and +0.055 once the group is admitted, a reduction of 33%; on long handling, +0.118 and +0.078, a reduction of 34%. Every rung sits outside its matched-dimension null, and across five split points the reduction ranges from 30% to 46%. Reopening is rare and its gains stand only 5.5 and 4.2 pooled standard deviations from their nulls, so it establishes the effect’s direction and not its size; and the reduction is largest on the target the opening group is most directly about, which is what Section 5 predicts. That is why we state the transferable claim as the gap’s existence rather than its magnitude.

8 One Thing We Could Not Establish

A third field, the knowledge-article reference, sits on the same Open row, is 100% populated, and costs nothing. Adding it

to the baseline raises AUC to 0.805 and drives the measured value of item identity to -0.003 . Whether that figure belongs in Table 1 depends on whether the field is available at creation, and we could not determine this.

We tried two structural tests and report both failures. A field that varies within an incident is certainly a per-event observation; the knowledge reference never varies. But neither does *Interaction ID*, which has 45,426 distinct values for 45,455 incidents and is plainly a creation-time key — so constancy is evidence of granularity, not of timing. We then tried exact identity with the closed record: the knowledge reference matches its closed-record column for 100.000000% of incidents. But restricted to the 42,151 incidents with exactly one related interaction, *Interaction ID* matches its own closed-record column for 99.997628% — one mismatch — against a pass mark of 100.000000%. The second test cannot separate its own counterexample either.

We therefore make no claim about this field. The -0.003 is also inside the dimensionality null of Section 4 and changes sign under penalty tuning and at other split points, so it is not resolvable on either ground. This is the paper’s own thesis turned on itself, and that reading is correct: our headline is contingent on a field-admission decision we cannot adjudicate, which is the exposure we argue every CMDB business case carries unstated. The transferable part is the negative — a field being free, populated and present on an opening record does not make it creation-time. The collection ships an interaction detail file we did not obtain; if it records this




Figure 3: (a) The shrinkage is graded in the opening group’s resolution. (b) How much of the item column’s contribution a partial CMDB recovers, averaged over five temporal split points; the band is the across-split range.

field at creation the question is answerable, and our inability to settle it is a limit of the three files we hold, not of the export.

9 Estate Concentration

Required CMDB scope is bounded by a property of the estate that needs no model. Here the top 8 items generate 30.2% of incidents, the top 64 generate 70.5%, and the top 128 — 5.0% of the training vocabulary — generate 82.0%. An organisation can compute this from a frequency count before funding anything.

Identifying only those items recovers most of the measured value. Averaged over five temporal split points, the top 8 recover 56% [53, 58] of the +0.103, the top 64 recover 88% [82, 92], and the top 128 recover 93% [91, 95] (Figure 3b). We report the mean and range rather than one split’s curve because a single curve is not monotone at this resolution: at $k = 32$ the across-split spread is 9 points, so adjacent values of k cannot be read against each other. The same holds with the group removed from the baseline — 89% at $k = 64$ — so this is a property of the estate’s concentration and not of our field-admission decision. We make no claim about what an

arbitrary coverage level recovers; an earlier draft did, from a comparison in which one of three rules had no observations where the claim was made.

10 Limitations

One organisation — of necessity, not choice: no other public ITSM log records the affected item (Section 2) — one platform, one task, six months. The data are from 2013 and 2014 and from a single on-premises ITSM product; estate composition, intake practice and tooling have all moved since, and we make no claim that the magnitudes transfer to a cloud-native estate. The values +0.183 and +0.103 are properties of this estate; the transferable claims are that the gap exists, that it is large, that it is graded in the free field’s resolution, and that the mechanism is the item column proxying for the opening group. Reassignment is a proxy for misrouting and also fires on legitimate escalation. The within-item shuffle preserves item membership exactly but not any other correlate of item identity, so it bounds the proxy share rather than isolating it perfectly. The opening group’s distribution drifts substantially between our training and test halves, which is realistic for a temporal split but means the group-aware baseline is being asked to generalise across a shift.

11 Conclusion

Knowing which configuration item an incident concerns is worth +0.183 or +0.103 AUC for predicting reassignment, depending on whether the baseline includes one free field: which group logged the ticket. Stated as detection at a fixed review capacity, the same choice moves the credited benefit several fold. The difference is an overlap: a model given the item has already been told most of what the group would say. A business case built on the first figure is not measuring a CMDB. It is measuring a CMDB plus a field the organisation was already recording for nothing.

Use of AI Systems

An AI assistant (Anthropic’s Claude) was used to implement and refactor analysis code, to construct adversarial challenges to intermediate claims — several withdrawals above originated that way — and to draft and edit prose the author reviewed and verified. All analyses were specified, run and checked by the author, who takes full responsibility for the content, including every numeric claim and every reference. The repository carries the working notes.

Acknowledgements

This work benefited from the review and domain challenge of a practitioner in IT service management; several analyses were withdrawn because of objections raised there.

Reproducibility

The dataset is public (van Dongen 2014). Code, figures and a checker comparing every quantity here against a generated result file: <https://github.com/sudhirdixit1/cmdb-routing-queue-baseline>.

References

- Amaral, C.; Fantinato, M.; and Peres, S. M. 2018. Incident Management Process Enriched Event Log. UCI Machine Learning Repository. Extracted from the audit system of a ServiceNow instance; 141,712 events over 24,918 incidents. CC BY 4.0.
- Bouthillier, X.; Delaunay, P.; Bronzi, M.; Trofimov, A.; Nichyporuk, B.; Szeto, J.; Mohammadi Sepahvand, N.; Raff, E.; Madan, K.; Voleti, V.; Ebrahimi Kahou, S.; Michalski, V.; Arbel, T.; Pal, C.; Varoquaux, G.; and Vincent, P. 2021. Accounting for Variance in Machine Learning Benchmarks. In *Proceedings of Machine Learning and Systems (MLSys)*.
- Chen, J.; He, X.; Lin, Q.; Zhang, H.; Hao, D.; Gao, F.; Xu, Z.; Dang, Y.; and Zhang, D. 2019. Continuous Incident Triage for Large-Scale Online Service Systems. In *Proceedings of the 34th IEEE/ACM International Conference on Automated Software Engineering (ASE)*.
- Covert, I.; Lundberg, S.; and Lee, S.-I. 2020. Understanding Global Feature Contributions With Additive Importance Measures. In *Advances in Neural Information Processing Systems 33 (NeurIPS)*.
- Ferrari Dacrema, M.; Cremonesi, P.; and Jannach, D. 2019. Are We Really Making Much Progress? A Worrying Analysis of Recent Neural Recommendation Approaches. In *Proceedings of the 13th ACM Conference on Recommender Systems (RecSys)*.
- Gundersen, O. E.; and Kjensmo, S. 2018. State of the Art: Reproducibility in Artificial Intelligence. In *Proceedings of the 32nd AAAI Conference on Artificial Intelligence*.
- Kapoor, S.; and Narayanan, A. 2023. Leakage and the Reproducibility Crisis in Machine-Learning-Based Science. *Patterns*, 4(9).
- Kaufman, S.; Rosset, S.; Perlich, C.; and Stitelman, O. 2012. Leakage in Data Mining: Formulation, Detection, and Avoidance. *ACM Transactions on Knowledge Discovery from Data*, 6(4).
- Lin, Q.; Hsieh, K.; Dang, Y.; Zhang, H.; Sui, K.; Xu, Y.; Lou, J.-G.; Li, C.; Wu, Y.; Yao, R.; Chintalapati, M.; and Zhang, D. 2018. Predicting Node Failure in Cloud Service Systems. In *Proceedings of the 26th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE)*.
- Marrone, M.; and Kolbe, L. M. 2011. Impact of IT Service Management Frameworks on the IT Organization: An Empirical Study on Benefits, Challenges and Processes. *Business & Information Systems Engineering*, 3(1): 5–18.
- Notaro, P.; Cardoso, J.; and Gerndt, M. 2021. A Survey of AIOps Methods for Failure Management. *ACM Transactions on Intelligent Systems and Technology*, 12(6).
- Remil, Y.; Bendimerad, A.; Mathonat, R.; and Kaytoue, M. 2024. AIOps Solutions for Incident Management: Technical Guidelines and A Comprehensive Literature Review. *arXiv preprint arXiv:2404.01363*.
- Rendle, S.; Zhang, L.; and Koren, Y. 2019. On the Difficulty of Evaluating Baselines: A Study on Recommender Systems. *arXiv preprint arXiv:1905.01395*.
- Steeman, W. 2013. BPI Challenge 2013, Incidents. 4TU.ResearchData. Volvo IT incident and problem management, VINST system.
- Suriadi, S.; Andrews, R.; ter Hofstede, A. H. M.; and Wynn, M. T. 2017. Event Log Imperfection Patterns for Process Mining: Towards a Systematic Approach to Cleaning Event Logs. *Information Systems*, 64: 132–150.
- Teinemaa, I.; Dumas, M.; La Rosa, M.; and Maggi, F. M. 2019. Outcome-Oriented Predictive Process Monitoring: Review and Benchmark. *ACM Transactions on Knowledge Discovery from Data*, 13(2): 17:1–17:57.
- van der Aalst, W. M. P. 2016. *Process Mining: Data Science in Action*. Springer, 2nd edition.
- van Dongen, B. F. 2014. BPI Challenge 2014. 4TU.ResearchData. Rabobank Group ICT, HP Service Manager; incident, change and interaction details.
- Verenich, I.; Dumas, M.; La Rosa, M.; Maggi, F. M.; and Teinemaa, I. 2019. Survey and Cross-Benchmark Comparison of Remaining Time Prediction Methods in Business Process Monitoring. *ACM Transactions on Intelligent Systems and Technology*.
- Weytjens, H.; and De Weerd, J. 2022. Creating Unbiased Public Benchmark Datasets with Data Leakage Prevention for Predictive Process Monitoring. In *Business Process Management Workshops*, volume 436 of *Lecture Notes in Business Information Processing*, 18–29. Springer.
- Williamson, B. D.; Gilbert, P. B.; Carone, M.; and Simon, N. 2021. Nonparametric Variable Importance Assessment Using Machine Learning Techniques. *Biometrics*, 77(1): 9–22.
- Zhang, L.; Jia, T.; Jia, M.; Wu, Y.; Liu, A.; Yang, Y.; Wu, Z.; Hu, X.; Yu, P. S.; and Li, Y. 2024. A Survey of AIOps for Failure Management in the Era of Large Language Models. *arXiv preprint arXiv:2406.11213*.